

Dynamic Programming

Optimal Substructure, Overlapping Subproblems, and Classical Applications

Asaduzzaman Herok

June 15, 2026

Part I — Foundations

- What is Dynamic Programming?
- Optimal Substructure
- Overlapping Subproblems
- Memoization (Top-Down)
- Tabulation (Bottom-Up)
- Space Optimisation
- Fibonacci — warm-up

Part II — Applications

- Coin Row / Coin Change
- 0-1 Knapsack + Memory Functions
- Optimal Binary Search Trees
- Warshall's Algorithm (Transitive Closure)
- Floyd's Algorithm (All-Pairs Shortest Paths)
- Complexity Summary Table

Reading

CLRS 4e — Chapter 14; Levitin — Chapter 8

Dynamic Programming — Core Concepts

Part I: Foundations

What is Dynamic Programming? I

Dynamic Programming (DP)

An algorithm design paradigm that solves complex problems by breaking them into **overlapping subproblems**, solving each subproblem **exactly once**, and **storing** the result for reuse.


History:

- Invented by **Richard Bellman** (1950s) for multistage decision optimization.
- “Programming” = *planning* (not coding).
- Now a general-purpose technique in CS.

Compared to Divide & Conquer:

- D&C: subproblems are **disjoint**
- DP: subproblems **overlap** — reuse is the key saving

What is Dynamic Programming? II



The Four-Step DP Method (CLRS)

- 1 **Characterise** the structure of an optimal solution.
— What does optimality look like?
- 2 **Recursively define** the value of an optimal solution.
— Write the recurrence $\text{OPT}(n) = f(\text{OPT}(\text{smaller}))$
- 3 **Compute** the optimal value, typically *bottom-up*.
— Fill a table, smallest subproblems first
- 4 **Construct** an optimal solution from the table.
— Backtrack through the table (if needed)

Key Idea

Steps 1–3 give the **value**; step 4 gives the **solution**.
Skip step 4 if only the value is needed.

1. Characterise Optimality



2. Write Recurrence



3. Compute Table



4. Reconstruct Solution

Optimal Substructure Property

A problem exhibits **optimal substructure** if an optimal solution to the problem contains optimal solutions to its subproblems. This is *Bellman's Principle of Optimality*.

Formally: If S^* is an optimal solution for problem P , and S^* is constructed from solutions $s_1, s_2, \dots$ to sub-problems $P_1, P_2, \dots$, then each s_i must be optimal for P_i .

Proof technique — “cut & paste”: Suppose s_i is not optimal. Replace it with s_i^* (the optimal). The resulting solution is strictly better — **contradiction**.

Holds for:

- Shortest paths (Bellman-Ford, Floyd)
- 0-1 Knapsack
- Optimal BST
- Coin change, Fibonacci

Fails for:

- Longest *simple* path in a graph (subpaths may share vertices)
- Unweighted longest common subsequence with constraints

Key Idea

Always verify optimality holds before applying DP!

Overlapping Subproblems

A problem has **overlapping subproblems** when the recursive algorithm visits the *same* subproblem multiple times. DP exploits this by computing each subproblem **only once**.


Why does this matter?

- Naive recursion: exponential time (recomputes!)
- DP (with memoisation): polynomial time

Contrast with D&C:

- Merge sort / quicksort: subproblems *disjoint*
- No benefit from storing results
- DP: subproblems *share* sub-subproblems

Fibonacci call tree for $F(5)$:



$F(3)$ computed **twice**, $F(2)$ **three times**!

Two DP Strategies: Memoisation vs. Tabulation I

Top-Down: Memoisation

- Write the natural recursion
 - Add a **cache** (hash map / array) initialised to `null`
 - Before computing, **check cache**
 - After computing, **store result**
 - Computes *only needed* subproblems
- + Intuitive, mirrors recursion
- + Lazy: skips unnecessary states
- Call-stack overhead (recursion)
- Harder to analyse precisely

Bottom-Up: Tabulation

- Identify ordering of subproblems
 - Fill a **table** from smallest to largest
 - No recursion — pure iteration
 - **All** subproblems computed
 - Straightforward space optimisation
- + No recursion overhead
- + Easy space reduction
- + Cache-friendly access pattern
- Must solve all subproblems

Two DP Strategies: Memoisation vs. Tabulation II

Key Distinction

Both achieve the same **asymptotic time complexity**. Tabulation is generally preferred for **space optimisation**. Memoisation is useful when the subproblem space is **sparse**.

Without memoization (exponential):

Naive Fibonacci

```
fib(n):  
    if n ≤ 1: return n  
    return fib(n-1) + fib(n-2)  
  
// Time:  $O(2^n)$  | recomputes  
// same calls
```

With memoization (linear):

Memoised Fibonacci

```
memo = {0: 0, 1: 1}  
  
fib(n):  
    if n in memo:  
        return memo[n]  
    memo[n] = fib(n-1) + fib(n-2)  
    return memo[n]  
  
// Time:  $O(n)$ , Space:  $O(n)$ 
```

General memoization

Template

```
// initialise table T to -1  
(null)  
T[0..n]  $\leftarrow$  -1  
  
solve(i):  
    if T[i]  $\geq$  0:  
        return T[i] // cached!  
  
    T[i]  $\leftarrow$  ...compute...  
    return T[i]
```

When to Use Memoisation

- Subproblem space is large but **sparse**
- Natural recursive formulation is clear
- Used in: Knapsack (memory functions), OBST

Tabulation (bottom-up):

Tabulated Fibonacci

```
fib_tab(n):  
    F[0] ← 0; F[1] ← 1  
    for i ← 2 to n:  
        F[i] ← F[i-1] + F[i-2]  
    return F[n]  
  
// Time: O(n), Space: O(n)
```

Execution trace ($n = 6$):

i	0	1	2	3	4	5	6
$F[i]$	0	1	1	2	3	5	8

General tabulation

Template

```
// 1. Set base cases  
F[0..base]  $\leftarrow$  known values  
  
// 2. Fill table in dependency  
order  
for i  $\leftarrow$  1 to n:  
    F[i]  $\leftarrow$  recurrence(F[...])  
  
// 3. Answer is in F[n]  
return F[n]
```

Key Insight

- The table fills in a topological order of subproblem dependencies.
- Every entry you compute depends only on **already-filled** entries.

Space Complexity Optimization I

Observation: Many DP recurrences only need a *small window* of previous values, not the full table.

General principle:

- If $F[i]$ depends only on $F[i - 1]$ and $F[i - 2]$: $\Rightarrow \mathcal{O}(1)$ space
- If $F[i][j]$ depends only on row $i - 1$: $\Rightarrow$ use two rows: $\mathcal{O}(n)$ instead of $\mathcal{O}(n^2)$
- Coin change (1D): $\mathcal{O}(n)$ trivially
- Knapsack: From $\mathcal{O}(nW)$ to $\mathcal{O}(W)$ (with careful traversal order — see later)
- OBST, Floyd: harder; often need $\mathcal{O}(n^2)$

Space-Optimized Fibonacci

```
fib_opt(n):  
    prev2 ← 0  
    prev1 ← 1  
    for i ← 2 to n:  
        curr ← prev1 + prev2  
        prev2 ← prev1  
        prev1 ← curr  
    return prev1  
  
// Time:  O(n), Space:  O(1)
```

Idea: Only keep last two values.

Trade-off Warning

Space-optimised DP often **loses the ability to backtrack** and reconstruct the solution. Keep full table if solution reconstruction is required.

Classical DP Problems— Deep Dives

Part II: Applications

Problem

Given n coins in a row with values $c_1, c_2, \dots, c_n$, pick up the maximum total value subject to: **no two adjacent coins may be picked.**

Optimal Substructure

Consider coin c_n (last):

- **Include** c_n : add c_n + best from coins $1 \dots n - 2$
- **Exclude** c_n : best from coins $1 \dots n - 1$

$$F(n) = \max\{c_n + F(n - 2), F(n - 1)\} \quad F(0) = 0, \quad F(1) = c_1$$

CoinRow(C[1..n])

```
F[0] ← 0; F[1] ← C[1]
for i ← 2 to n do
    F[i] ← max(C[i] + F[i-2], F[i-1])
return F[n]
```

Backtrack: Optimal = $\{c_1, c_4, c_6\} = \{5, 10, 2\}$

Complexity

Time: $\Theta(n)$

Space: $\Theta(n)$ (or $\mathcal{O}(1)$ with just 2 vars)

Example: Coins = $[5, 1, 2, 10, 6, 2]$

i	0	1	2	3	4	5	6
$C[i]$	—	5	1	2	10	6	2
$F[i]$	0	5	5	7	15	15	17

Calculation:

$$F[2] = \max(1 + 0, 5) = 5$$

$$F[3] = \max(2 + 5, 5) = 7$$

$$F[4] = \max(10 + 5, 7) = 15$$

$$F[5] = \max(6 + 7, 15) = 15$$

$$F[6] = \max(2 + 15, 15) = \mathbf{17}$$

Problem

Given denominations $d_1 < d_2 < \dots < d_m$ (with $d_1 = 1$) and amount n , find the **minimum number of coins** that sum to n .

Optimal Substructure:

The last coin used has denomination d_j . After using it, we need minimum coins for $n - d_j$.

- Try every valid denomination $d_j \leq n$
- Pick the one minimising $F(n - d_j) + 1$

$$F(n) = \min_{j: n \geq d_j} \{F(n - d_j)\} + 1$$
$$F(0) = 0$$

ChangeMaking($D[1..m]$, n)

```
F[0]  $\leftarrow$  0
for i  $\leftarrow$  1 to n do
    temp  $\leftarrow$   $\infty$ ; j  $\leftarrow$  1
    while j  $\leq$  m and i  $\geq$  D[j] do
        temp  $\leftarrow$  min(F[i - D[j]], temp)
        j  $\leftarrow$  j + 1
    F[i]  $\leftarrow$  temp + 1
return F[n]
```

Example: $n = 6$, denominations
 $= \{1, 3, 4\}$

n	0	1	2	3	4	5	6
$F[n]$	0	1	2	1	1	2	2

Key computations:

$$F[1] = F[0] + 1 = 1 \quad (d = 1)$$

$$F[3] = \min(F[2], F[0]) + 1 = 1 \quad (d = 3)$$

$$F[4] = \min(F[3], F[1], F[0]) + 1 = 1 \quad (d = 4)$$

$$F[6] = \min(F[5], F[3], F[2]) + 1 = 2 \quad (d = 3)$$

Solution: Two coins of denomination 3

Backtrack: by recording which d_j achieved minimum.

Complexity

Time: $\mathcal{O}(nm)$ Space: $\Theta(n)$

Variant 1 — Minimum Coins (done)

- Recurrence:
$$F(n) = \min_{d_j \leq n} \{F(n - d_j)\} + 1$$
- Unlimited coin supply
- Time: $\mathcal{O}(nm)$, Space: $\mathcal{O}(n)$

Variant 2 — Count of Ways

- How many ways to make change for n ?
- $\text{ways}(n) = \sum_{d_j \leq n} \text{ways}(n - d_j)$
- Same structure, addition instead of min

Variant 3 — Limited Supply

- At most k_j coins of denomination d_j
- Becomes a bounded knapsack variant
- Extra dimension for coin count

Variant 4 — Coin Row (no adjacency)

- As seen:
$$F(n) = \max(c_n + F(n-2), F(n-1))$$
- Fixed arrangement, each coin used at most once

Greedy vs. DP:

Key Idea

Greedy (“always use largest coin”) works for standard US/EU denominations, but **fails** in general. Example:

Denominations: $\{1, 3, 4\}$, Amount = 6

Greedy: $4 + 1 + 1 = 3$ coins

DP: $3 + 3 = 2$ coins ✓

DP is the correct general approach.

0-1 Knapsack

Given n items with weights $w_1, \dots, w_n$ and values $v_1, \dots, v_n$, and a knapsack of capacity W : select a subset of items to **maximise total value** subject to **total weight** $\leq W$. Each item is either included (1) or not (0).

Optimal Substructure


Let $F(i, j) = \text{max value from items } 1 \dots i \text{ with capacity } j$.

Two cases for item i :

- **Leave item i :** $F(i - 1, j)$
- **Take item i (if $j \geq w_i$):** $v_i + F(i - 1, j - w_i)$

$$F(i, j) = \begin{cases} \max\{F(i-1, j), v_i + F(i-1, j - w_i)\} & j \geq w_i \\ F(i-1, j) & j < w_i \end{cases}$$
$$F(0, j) = 0 \quad \forall j; \quad F(i, 0) = 0 \quad \forall i$$

State space interpretation:



Fill order: Row by row (or column by column). Each cell depends on the row **above** it.

Complexity

Time: $\Theta(nW)$ Space: $\Theta(nW)$

Note: *pseudo-polynomial* — W is the numeric value, not bit-length

0-1 Knapsack — Worked Example I — Capacity $W = 5$

Items:

Item	1	2	3	4
Weight	2	1	3	2
Value	\$12	\$10	\$20	\$15

DP Table $F(i, j)$:

$i \setminus j$	0	1	2	3	4	5
0 (base)	0	0	0	0	0	0
1 ($w=2, v=12$)	0	0	12	12	12	12
2 ($w=1, v=10$)	0	10	12	22	22	22
3 ($w=3, v=20$)	0	10	12	22	30	32
4 ($w=2, v=15$)	0	10	15	25	30	37

Optimal value: $F(4, 5) = \$37$

Backtracking to find composition:

- $F(4, 5) > F(3, 5) \Rightarrow$ include item 4 ($w = 2$); go to $F(3, 3)$
- $F(3, 3) = F(2, 3) \Rightarrow$ skip item 3
- $F(2, 3) > F(1, 3) \Rightarrow$ include item 2 ($w = 1$); go to $F(1, 2)$
- $F(1, 2) > F(0, 2) \Rightarrow$ include item 1 ($w = 2$)

$\Rightarrow$ **Optimal set:** {item 1, item 2, item 4}

Motivation

Bottom-up computes **all** $n \times W$ cells. For some instances, only a **small subset** of cells is needed. Memory functions (memoised top-down) compute **only necessary** subproblems.

Memory Functions: Top-Down DP for Knapsack II

Algorithm:

MFKnapsack(i, j)

```
// F init: row 0 and col 0 = 0, rest = -1
if F[i, j] ≥ 0:
    return F[i, j] // cached
if j < Weights[i]:
    val ← MFKnapsack(i-1, j)
else:
    val ← max(MFKnapsack(i-1, j), Values[i] + MFKnapsack(i-1, j - Weights[i]))
F[i, j] ← val
return F[i, j]
```

Memory Functions: Top-Down DP for Knapsack III

Computed cells for the worked example:

$i \backslash j$	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	12	12	12	12
2	0	—	12	22	—	22
3	0	—	—	22	—	32
4	0	—	—	—	—	37

“—” = never computed; blue = computed.

Analysis

Only **11 of 20** non-trivial cells computed. Efficiency class: same $\mathcal{O}(nW)$ worst case, but can be significantly faster in practice.

Knapsack: Space Optimisation to $\mathcal{O}(W)$ I

Key observation: $F(i, j)$ depends only on **row** $i - 1$ of the table. $\Rightarrow$ We only need **one array of size** $W + 1$, updated in-place.

Critical: Must iterate j from **right to left** (W down to w_i) to avoid using the same item twice.

Space-Optimised Knapsack

```
F[0..W]  $\leftarrow$  0 // single array
for i  $\leftarrow$  1 to n do
    for j  $\leftarrow$  W downto w[i] do
        F[j]  $\leftarrow$  max(F[j], v[i] + F[j - w[i]])
return F[W]
```

Knapsack: Space Optimisation to $\mathcal{O}(W)$ II

Complexity

Time: $\Theta(nW)$ Space: $\Theta(W)$

Why right-to-left?

Before item i :

0	10	12	22	22	22
---	----	----	----	----	----

After item i ($w=2, v=15$):

0	10	15	25	30	37
---	----	----	----	----	----

update order



Warning

Iterating **left-to-right** would allow an item to be selected **multiple times** — this solves the unbounded knapsack instead!

Trade-off: Cannot reconstruct optimal subset directly; need additional bookkeeping if the subset is required.

Optimal Binary Search Tree — Motivation I

Problem

Given n keys $a_1 < a_2 < \dots < a_n$ with search probabilities $p_1, p_2, \dots, p_n$ ($\sum p_i = 1$), build the BST that **minimises the expected number of comparisons** in a successful search.

Expected comparisons in tree T :

$$E[T] = \sum_{i=1}^n p_i \cdot \text{depth}(a_i)$$

where depth of root = 1.

Why not just build any BST?

- A balanced BST minimises *worst-case* depth
- OBST minimises *expected* comparisons
- High-probability keys should be near root


Key Idea

Total BSTs with n keys = Catalan number $C_n \sim \frac{4^n}{n^{1.5}}$

Exhaustive search is completely infeasible.

Example: Keys A, B, C, D with probs 0.1, 0.2, 0.4, 0.3

Optimal Binary Search Tree — Motivation III



$$E = 0.1 \cdot 3 + 0.2 \cdot 2 + 0.4 \cdot 1 + 0.3 \cdot 2 = 1.7$$

This is the **optimal** BST for this example.

Key C (highest probability 0.4) is at the root.

Optimal BST — Recurrence and Algorithm I

Let $C(i, j) = \min$ expected comparisons in optimal BST for keys $a_i, \dots, a_j$.

Choose root a_k ($i \leq k \leq j$):

- Left subtree: optimal BST for $a_i, \dots, a_{k-1}$
- Right subtree: optimal BST for $a_{k+1}, \dots, a_j$
- When entire subtree becomes a child, depths +1

$$C(i, j) = \min_{i \leq k \leq j} \{C(i, k-1) + C(k+1, j)\} + \sum_{s=i}^j p_s$$


$$C(i, i-1) = 0 \text{ (empty tree); } C(i, i) = p_i$$

OptimalBST(P[1..n]) — sketch

```
for d ← 1 to n-1 do // diagonal
  for i ← 1 to n-d do
    j ← i + d
    for k ← i to j do
      find min C[i,k-1] + C[k+1,j]
    C[i,j] ← minval + sum(P[i..j])
return C[1,n], R
```

Table filling order (diagonal by diagonal):

Optimal BST — Recurrence and Algorithm III



Result table for A,B,C,D:

Optimal BST — Recurrence and Algorithm IV

$C(i, j)$	$j=1$	$j=2$	$j=3$	$j=4$
$i=1$	0.1	0.4	1.1	1.7
$i=2$	—	0.2	0.8	1.4
$i=3$	—	—	0.4	1.0
$i=4$	—	—	—	0.3


Complexity

Time: $\mathcal{O}(n^3)$ (naive) or $\mathcal{O}(n^2)$ (Knuth) Space: $\mathcal{O}(n^2)$

Optimal BST — Solution Reconstruction I

Root Table $R(i, j)$: Record which k achieved the minimum in $C(i, j)$.

$R(i, j)$	$j=1$	$j=2$	$j=3$	$j=4$
$i=1$	1	2	3	3
$i=2$	—	2	3	3
$i=3$	—	—	3	3
$i=4$	—	—	—	4



Reconstruction procedure:

- 1 $R(1, 4) = 3$: root is key **C**
- 2 Left subtree: $R(1, 2) = 2$: root is **B**
- 3 $R(1, 1) = 1$: left child of B is **A**
- 4 Right subtree: $R(4, 4) = 4$: root is **D**

Verification:

$$\begin{aligned} E &= 0.1 \cdot 3 + 0.2 \cdot 2 + 0.4 \cdot 1 + 0.3 \cdot 2 \\ &= 0.3 + 0.4 + 0.4 + 0.6 = \mathbf{1.7} \end{aligned}$$

Compare to balanced tree $E = 2.1$ — OBST is better!

Key Idea

Knuth's improvement: Root table entries are monotone along rows/cols, restricting k search to $[R(i, j-1), R(i+1, j)]$ — achieves $\mathcal{O}(n^2)$.